

# Operating systems and concurrency B07

David Kendall

Northumbria University

- Semaphores can be used to solve a number of classical synchronisation problems
- We will consider:
  - Producer/consumer problem
  - Readers/writers problem

# Producer/consumer problem

- Very often in OS and concurrent applications programs, we have one or more tasks that **produce** data that must be used (**consumed**) by some other task(s).
- The rate at which data is produced may be, occasionally, greater than the rate at which data can be consumed
- A **buffer** can be useful to smooth out the differences in the rates of production and consumption

# Producer/consumer problem

## Real-world analogy

Imagine two people washing up. One person (the **producer**) washes the dishes and puts them in the dish rack (the **buffer**). The other person (the **consumer**) takes the dishes from the dish rack and dries them. If the dish rack fills up, the washer has to wait until the drier takes a dish from the rack. If the rack is empty, the drier has to wait for the washer to wash another dish and put it in the rack. (from [Goetz et al., 2006])

- Our problem is to implement the dish rack . . .
- . . . and to ensure that the washer-up and drier use it properly

# Naive circular buffer (.h)

```
#ifndef __BUFFER_H
#define __BUFFER_H

enum {
    BUF_SIZE = 6UL
};

typedef struct message {
    unsigned int data;
} message_t;

void putBuffer(message_t const * const);
void getBuffer(message_t * const);

#endif
```

# Naive circular buffer (.c)

```
#include <buffer.h>

static message_t buffer[BUF_SIZE];
static unsigned int front = 0;
static unsigned int back = 0;

void putBuffer(message_t const * const msg) {
    buffer[back] = *msg;
    back = (back + 1) % BUF_SIZE;
}

void getBuffer(message_t * const msg) {
    *msg = buffer[front];
    front = (front + 1) % BUF_SIZE;
}
```

# Naive circular buffer (Use)

```
/* producer */  
#include <buffer.h>
```

```
message_t msg;
```

```
...
```

```
msg.data = 27;  
putBuffer(&msg);
```

```
/* consumer */  
#include <buffer.h>
```

```
message_t msg;
```

```
...
```

```
getBuffer(&msg);  
lcdWrite('%u', msg.data);
```

# Problems with a naive buffer

- Interference between producer(s) and consumer(s)
  - Imagine two producers (P1 and P2) concurrently executing `putBuffer`: P2 does `buffer[back] = ...` and is then descheduled; P1 starts and finishes `putBuffer(...)`; P2 finishes `putBuffer(...)`.
  - What has gone wrong? Draw the state of the buffer.
- Attempt to put data into a full buffer
  - No room on the dish rack – **must wait**.
- Attempt to get data from an empty buffer
  - No dishes to dry – **must wait**.



# Elements of a solution

- Enforce mutual exclusion to avoid interference
  - Semaphore `bufMutex` initialized to the value 1
- Enforce producer wait if no buffer slots are empty
  - Semaphore `emptySlot` initialized to the value `BUF_SIZE`
- Enforce consumer wait if no buffer slots are full
  - Semaphore `fullSlot` initialized to the value 0

# Structure of producer

## Pseudo-code

```
while (true)

    // produce an item

    wait (emptySlot);
    wait (bufMutex);

    // add the item to the buffer

    post (bufMutex);
    post (fullSlot);

}
```

# Structure of consumer

Pseudo-code

```
while (true) {  
    wait (fullSlot);  
    wait (bufMutex);  
  
    // remove an item from buffer  
  
    post (bufMutex);  
    post (emptySlot);  
  
    // consume the item  
  
}
```

- Multiple tasks require access to a shared data structure, database or filesystem
- Some tasks only need to **read** the data (readers)
- Other tasks only need to **write** the data (writers)
- We need to avoid interference (how might this occur?)
- Full mutual exclusion may be inefficient (why?)
- So we require:
  - 1 Any number of readers can be allowed to read simultaneously
  - 2 A writer must have exclusive access (ie no other writers and no readers can access the data at the same time as the writer)

# Elements of a solution

- Ensure mutually exclusive access to the data when writing
  - Semaphore `writeMutex` initialised to 1
- Keep a count of the number of readers
  - unsigned int `nReaders` initialised to 0
- Ensure mutually exclusive access to the readers count
  - Semaphore `nReadersMutex` initialised to 1

# Structure of writer

## Pseudo-code

```
while (true) {  
    wait(writeMutex);  
  
    /* write the data */  
  
    post(writeMutex);  
  
    /* do non-critical stuff */  
}
```

- Protocol for a writer is very simple...
- ...it needs exclusive access to the data

# Structure of a reader

## Pseudo-code

```
while (true) {  
    wait(nReadersMutex);  
    nReaders += 1;  
    if (nReaders == 1) {  
        wait(writeMutex);  
    }  
    post(nReadersMutex);  
  
    /* read the data */  
  
    wait(nReadersMutex);  
    nReaders -= 1;  
    if (nReaders == 0) {  
        post(writeMutex);  
    }  
    post(nReadersMutex);  
  
    /* do non-critical stuff */  
}
```

- We keep track of the number of readers
- The first reader to arrive needs to ensure that there are no writers  
*wait(writeMutex)*
- The last reader to finish should release any waiting writers *post(writeMutex)*
- Readers that arrive while the first reader is waiting for a writer to finish are suspended by the first  
*wait(nReadersMutex)*

# Problem with this solution

- There's a problem with this solution that we'll consider in a later session.



# Summary

- Semaphores allow us to solve a variety of synchronisation problems
- Care is required to avoid a number of problems ...
- ... to be considered later

# Acknowledgements

- Silberschatz, Galvin, Gagne, *Operating System Concepts*, John Wiley, 2013
- B. Goetz with T. Peierls, J. Bloch, J. Bowbeer, D. Holmes, and D. Lea, *Java Concurrency in Practice*, Addison Wesley, 2006